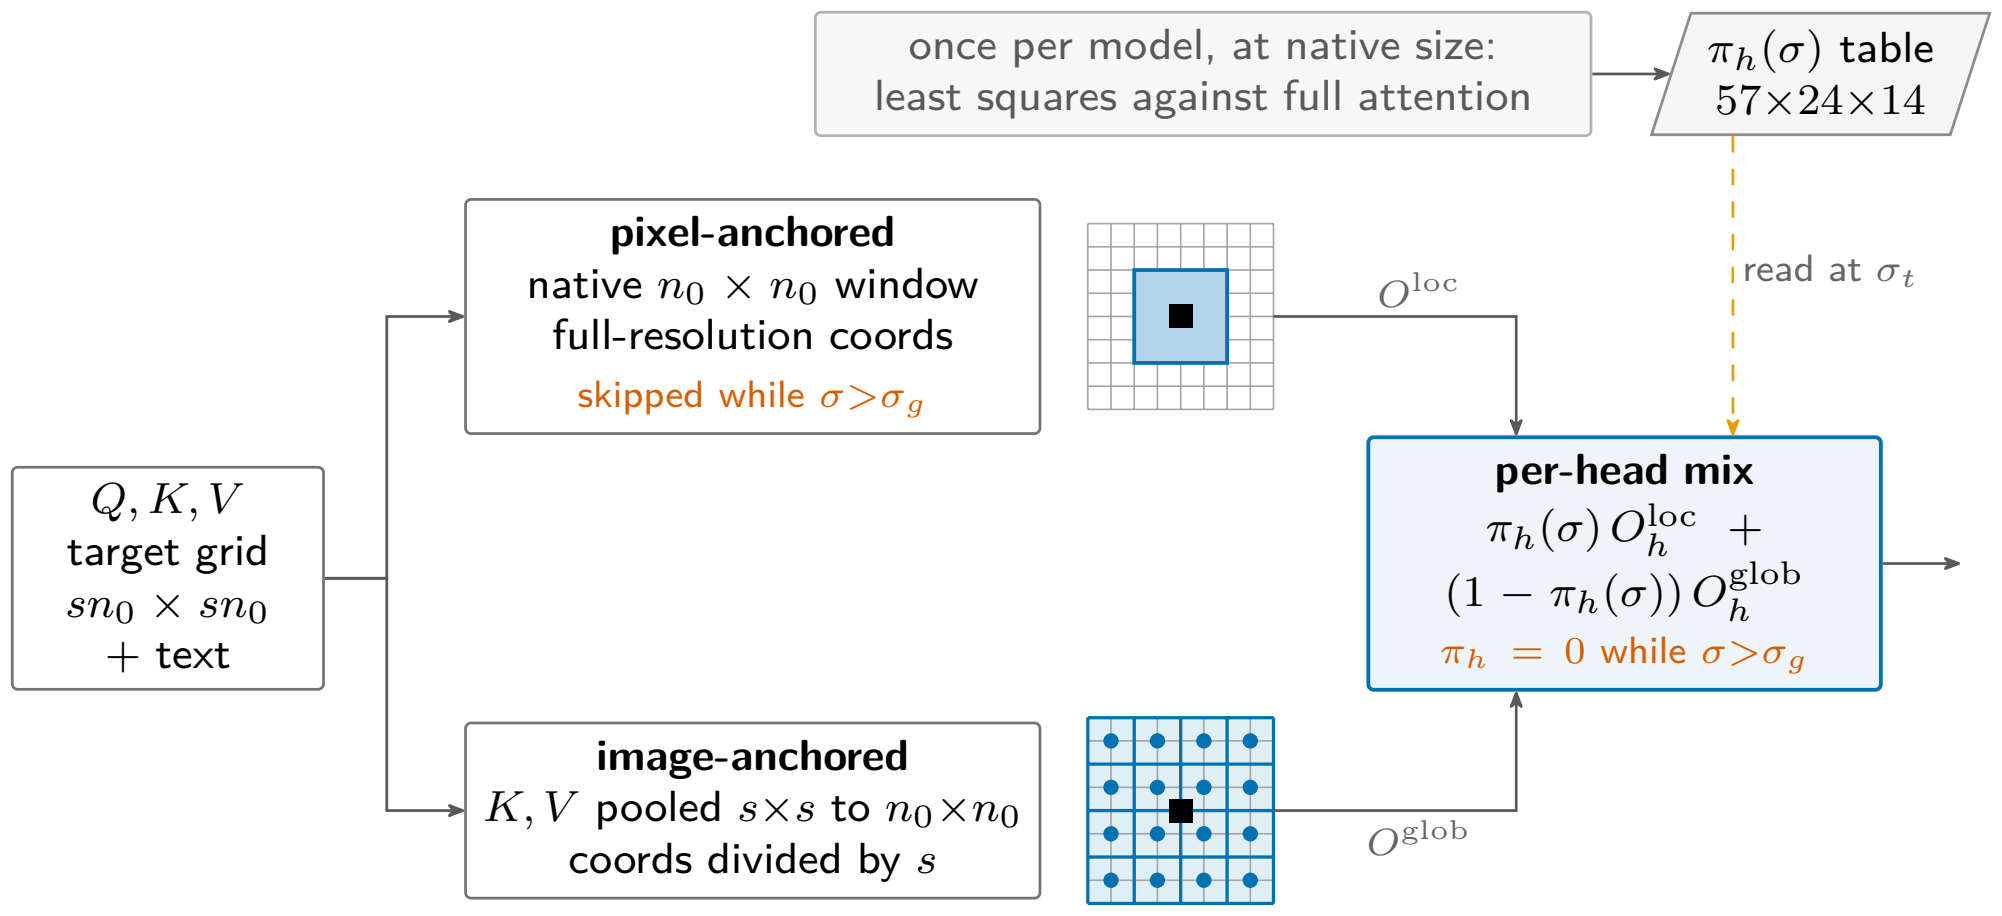
once per model, at native size:
least squares against full attention
$\pi_h(\sigma)$ table
$57\times24\times14$
read at σ_t
Q, K, V
target grid
$sn_0 \times sn_0$
+ text
pixel-anchored
native $n_0 \times n_0$ window
full-resolution coords
skipped while $\sigma>\sigma_g$
O^{loc}
per-head mix
$\pi_h(\sigma)\, O_h^{\text{loc}} +$
$(1-\pi_h(\sigma))\, O_h^{\text{glob}}$
$\pi_h = 0$ while $\sigma>\sigma_g$
image-anchored
K, V pooled $s\times s$ to $n_0\times n_0$
coords divided by s
O^{glob}